

Hardware Root Of Trust

Chapter 5

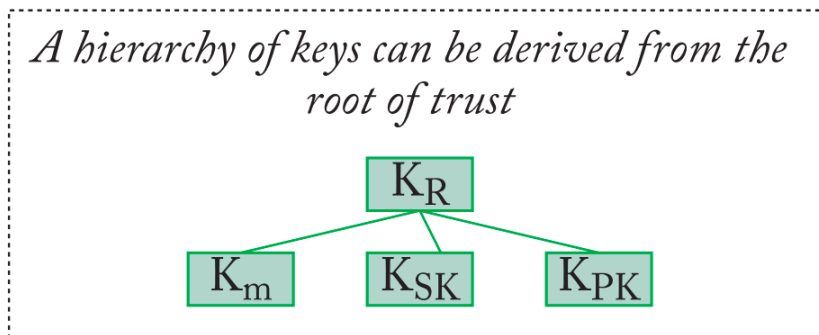
- [1] J. Szefer, “Principles of secure processor architecture design,” *Synth. Lect. Comput. Archit.*, vol. 13, no. 3, pp. 1–173, 2018.

The Root of Trust

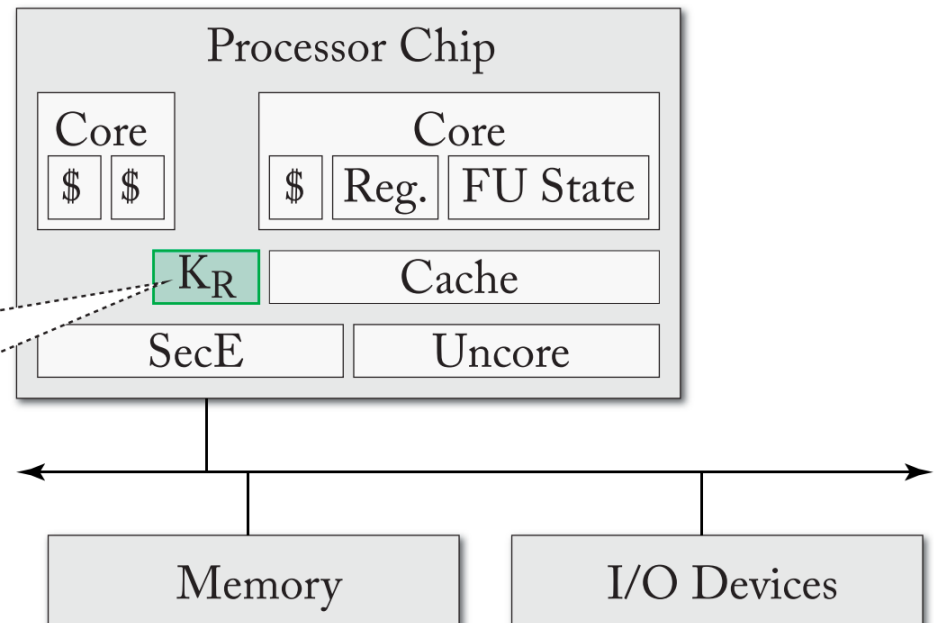
- ▣ Measure the chain of trust
 - ◆ Trusted authenticated boot, remote attestation, and data sealing
 - ◆ Runtime attestation and continuous attestation
- ▣ Secure processor should ensure for TEE
 - ◆ **Authentication**: the processor should authenticate himself (to guarantee that the information coming from it can be trusted).
 - E.g., deriving asymmetric crypto keys from a manufacture dependent unique secret key
 - ◆ **Confidentiality**: data going out of processor should be protected.
 - E.g., For memory using a key created at boot time (ephemeral) for persistent keys are persistent and generated (in a systematic way)
 - ◆ **Integrity**: additional keys are required for integrity checking. Cryptographic hashing (one-way) from secret key to prevent tampering (e.g., replays)

The Processor Key (The root of trust)

- **K_R** is a unique key per chip: is the root of trust. Others derived from that with one-way hashing
- **Never should leave the chip boundary**
- Created at manufacturing time (e.g., eFUSE)
- Manufacturer **shouldn't** hold a copy of the key
- Use PUF?, more costly to guarantee reliability/stability

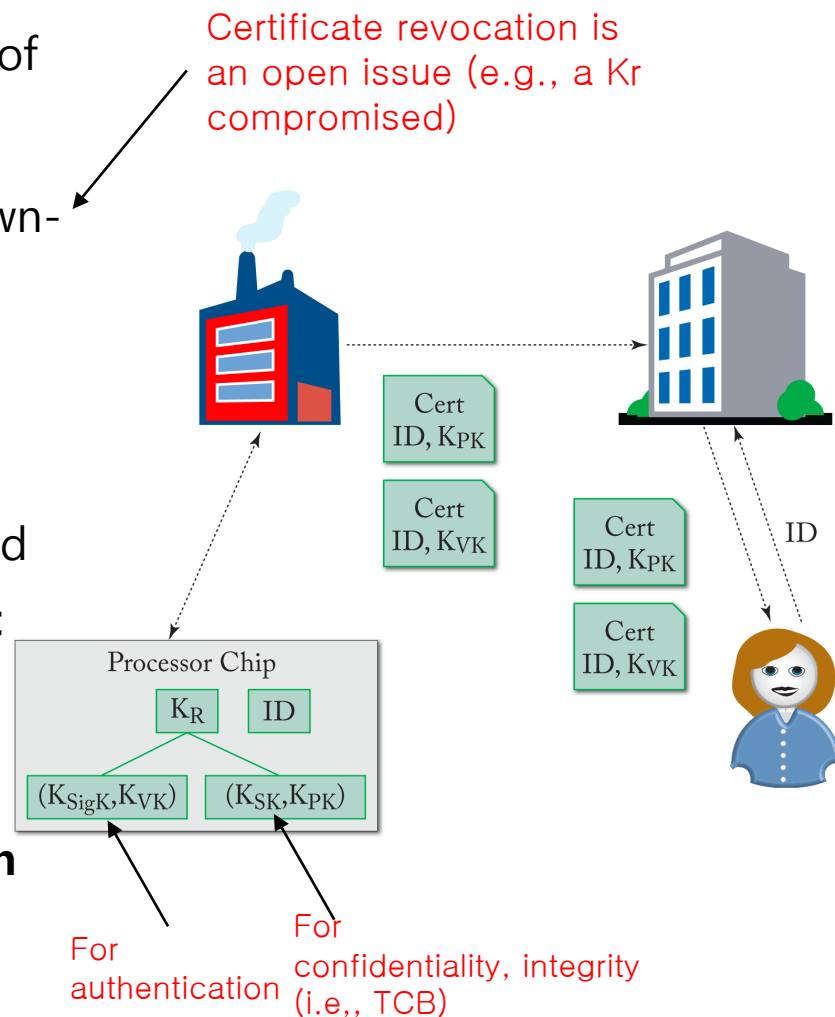


Keys belong to TC, if TEE need keys, they will be generated from them



PKI and Secure Processors

- **How to tell if a processor is who it says he is?**
- PKI is required to distributed the certificate of the public key of the processors
- The manufacturer should keep a list of known-good keys and the associated processors ID
- K_r is used to derive different pairs for **encryption** (K_{sk} , K_{pk}) or **signing** (K_{SigK} , K_{V_k})
- ID, public (K_{pk}) and verification (K_{vk}) keys, and corresponding certificates, are **generated at manufacture** and given to the certificate authority
- Given some ID of a processor, **the users can obtain from CA the certificates for that processor's keys**



Aside: RSA and Asymmetric Cryptography

1. Generate **two large** random primes, p and q , of approximately equal size such that their product $n=pq$ is of the required bit length, e.g., 1024 bits.
 2. Compute $n=pq$ and $\phi=(p-1)(q-1)$
 3. Choose an integer e , $1 < e < \phi$, such that $\gcd(e, \phi) = 1$
 4. Compute the secret exponent d , $1 < d < \phi$, such that $e \cdot d \equiv 1 \pmod{\phi}$
 5. The public key is (e, n) and the private key (d, n) . Keep other values secret.
-
- ▣ n is known as the *modulus*.
 - ▣ e is known as the **public exponent** or **encryption exponent** or just the *exponent*.
 - ▣ d is known as the **secret exponent** or **decryption exponent**.

Aside: RSA Example

- Choose $p = 3$ and $q = 11$
- Compute $n = p * q = 3 * 11 = 33$
- Compute $\phi(n) = (p - 1) * (q - 1) = 2 * 10 = 20$
- Choose e such that $1 < e < \phi(n)$ and e and $\phi(n)$ are coprime.
Let $e = 7$
- Compute a value for d such that $(d * e) \% \phi(n) = 1$.
One solution is $d = 3 [(3 * 7) \% 20 = 1]$
- Public key is $(e, n) \Rightarrow (7, 33)$
- Private key is $(d, n) \Rightarrow (3, 33)$
- The encryption of $m = 2$ is $c = 2^7 \% 33 = 29$
- The decryption of $c = 29$ is $m = 29^3 \% 33 = 2$

Aside: Computational Cost of RSA

- Public key operations $O(n^2)$, Private Key operations $O(n^3)$


- Very costly when size of the key increases

- Many $A*B \bmod N$ operations

- No support in current processors but SIMD (e.g., AVX) might help
- External hardware support; e.g., PCI Cards, FPGA, ...

- ECDSA [Elliptic Curve Cryptography] is replacing RSA everywhere (**224-bit ECDSA** is equivalent, from security standpoint, to **2048-bit RSA**)

- Post quantum: **Dilithium, Rainbow and Falcon**



Access to the Root Of Trust

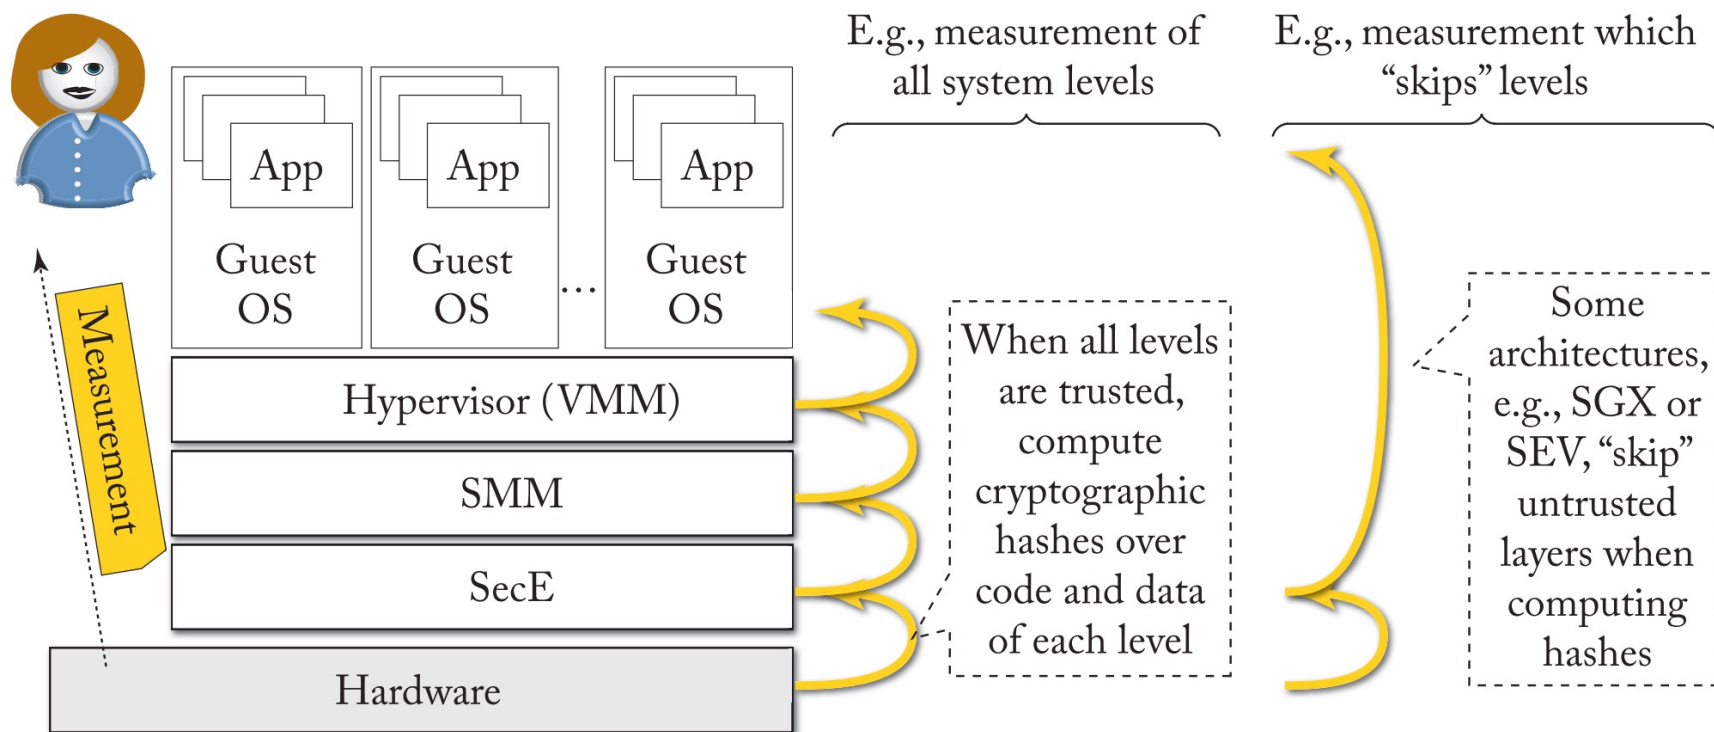
- ▣ Although complex physical attacks and/or rogue manufacturer can compromise the Root of Trust, the most likely scenario is **bugs in TCB**
 - ◆ Low level code (e.g., UEFI) should be able to access secret keys (e.g., to sign digital messages generated by the processor)
 - ◆ Code is kept in flash memory
 - ◆ Bugs in that code can lead to exploits that can leak the secret key

Chain of Trust and Measurements

- ▣ When run TEE, is crucial to be correctly configured
 - ◆ A trusted small BIOS loaded at system boot (first thing to load!)
 - ◆ This piece can "**measure**" next piece to be loaded (e.g., firmware)
 - ◆ Firmware "**measure**" the OS that starts next
 - ◆ ...
- ▣ **Measure** == act of generating a cryptographic hash of the code (and data) before it is executed
 - ◆ If **measurement** != **expectations**, **tampering is detected**
- ▣ Measurements can be chained (or completed) from multiple components.
 - ◆ Each measurement is "**extend**" with the next one in the chain (coined by TPM)
 - ◆ This idea is used by **Trusted Platform Module** (TPM): low cost dedicated to providing root of trust. Protects against software attacks (mainly **rootkits**), secure **disk cyphering**, machine **authentication** (e.g., windows hello)

Chain of Trust and Measurements

- ▣ Chains of trust can include untrusted levels
- ▣ **Usually**, only trusted components are measured
 - ◆ **Only if software is part of TCB and TEE is measured**



Trusted and Authenticated Boot

- ▣ **Trusted boot** means system might boots only if all measurements are correct
- ▣ **Authenticated boot** means users can access to all measurements taken during boot (but always boot if aren't correct)
- ▣ In **trusted boot** If measurements aren't correct no standardized actions
 - ◆ A system can stop executing or secret data won't be decrypted
 - ◆ But what if the system is part of a **critical infrastructure**?
 - ◆ But what is my laptop and I want to prevent Linux + passwd cracker to unencrypt my disk content?

Measurement Validation and Remote Attestation

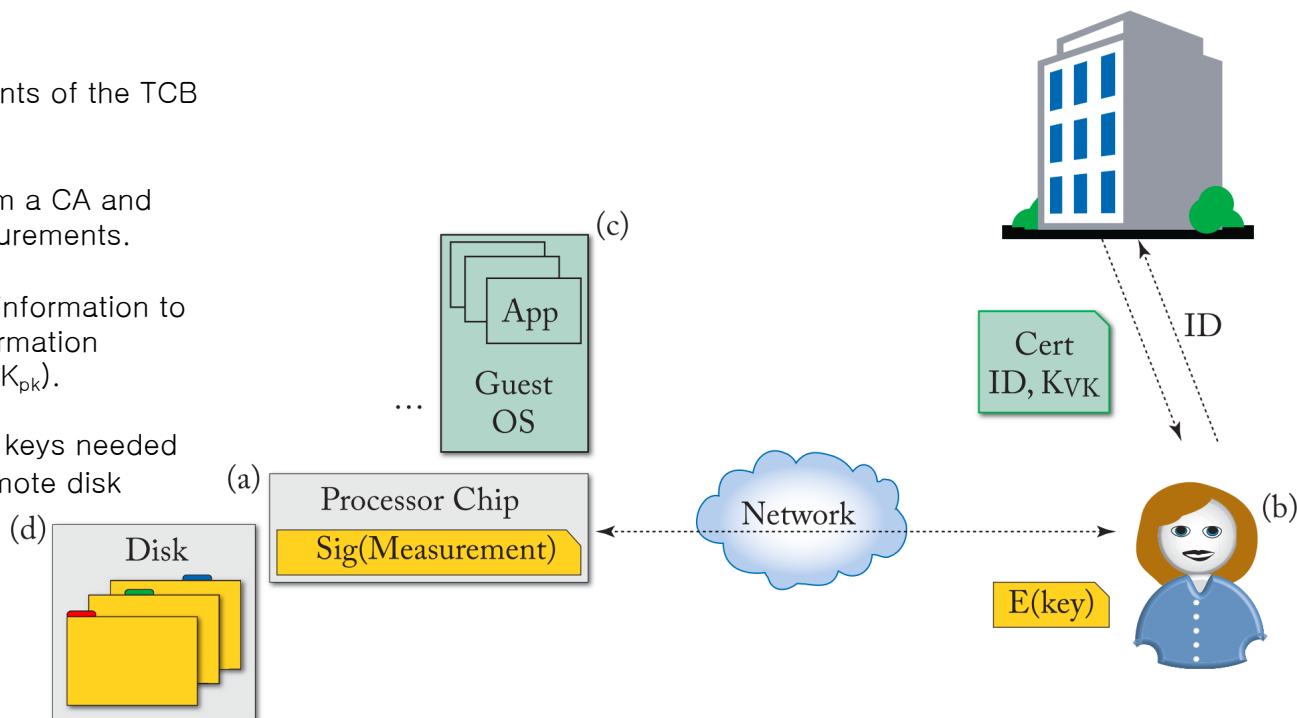
- Main method: check measurements against a list of **good values**
- But any change in the boot (e.g., change in the version of the loader) might completely change the resulting hash
 - ◆ **Can't have all combinations** of system software locally "tagged": Use PKI for key distribution

(a) Processor generate measurements of the TCB (and TEE) signed with K_r .

(b) User can obtain certificates from a CA and validate the signatures of the measurements.


(c) User can send some sensitive information to the TEE, by sending encrypted information (and decryption key protected with K_{pk}).

(d) The user can also send to TEE keys needed for decryption of data stored on remote disk



Example: Remote Attestation with Intel SGX

- Allows a hardware entity or combination of hardware and software to gain the trust of a remote provider (also known as a relying party) or manufacturer.



- Intel ME Core iX 11xxx<
- Intel SPS All Xeon >D >2020
- (Server Platform Services)

Sealing

- Act of **encrypting** some data with a key derived from a measurement.
- Unsealing, is the act of decrypting data with a key derived from a measurement.
- Tie the data to a specific hardware and software configuration.
 - ◆ For example, if **hard drive data is sealed at shutdown (local or remote)**
 - ◆ Next time the system boots up, if there is any change to the measurements (e.g., any of the trusted software was somehow **changed**), then the proper decryption key **cannot be derived**, and data cannot be decrypted.
- This prevents malicious users from accessing data if they changed the software while the system was turned off (e.g., **replacing the boot disk**)

Time-to-check to time-to-use attacks

- ▣ Measurement technique does not say anything about the code **after it was measured**.
 - ◆ It only is based on raw measurements of static code as it was right before the piece of software executed.
 - ◆ This leads to well-known class of Time-of-Check to Time-of-Use (TOC-TOU) problems
- ▣ The system may be **correct at time t_0** , but it may become **compromised at time t_1**
 - ◆ if user **asks for the values of measurements at some later time t_2** , he or she will get the hash value that states that system was okay at t_0 , but no information about state at t_2 is gained
- ▣ One solution to this problem is to perform **continuous attestation** at system runtime.

Runtime Attestation and Continuous Monitoring of TCB and TEEs



- ▣ **Monitoring** of the execution of the system via co-processor
 - ◆ Validates continuously the execution path (combined with PMU)
- ▣ Observing the execution pattern (Basic Block size, branch predictor accuracy, performance metric ...) we can infer when a program deviates from expected behavior
 - ◆ Mostly applied to untrusted component (Use in TCB not explored)
- ▣ Using Control Flow Graph of the program (from the source code of the program) to check if there are deviations (valid only for software, not clear for hardware)
- ▣ Using PMU to detect anomalies can be **easily defeated** by hiding the attack "in the noise" (e.g., by carefully modifying slightly the normal execution characteristics)

Physically Unclonable Functions (PUF) and Root of Trust

- PUFs are **impossible to replicate** (even for the manufacturer) and provides a good fingerprint for the systems
 - ◆ Turns process variation problem into a virtue
- Can permanently **bind hardware-software** by condition TEE execution path (i.e., not only in measurements)
 - ◆ The PUF might generate a stream of outputs for the TEE
 - E.g., Generate a seed for a PRNG
 - E.g., Use for some form of indirection (DRAM PUF)


Aside: Physically Unclonable functions (PUF)

- Circuits that can't be replicated because depends upon small variabilities during the manufacturing process (die-to-die variability)
 - ◆ Mainly random processes across wafer



Arbiter PUF:

Which path is faster determining the output for the challenge (X)

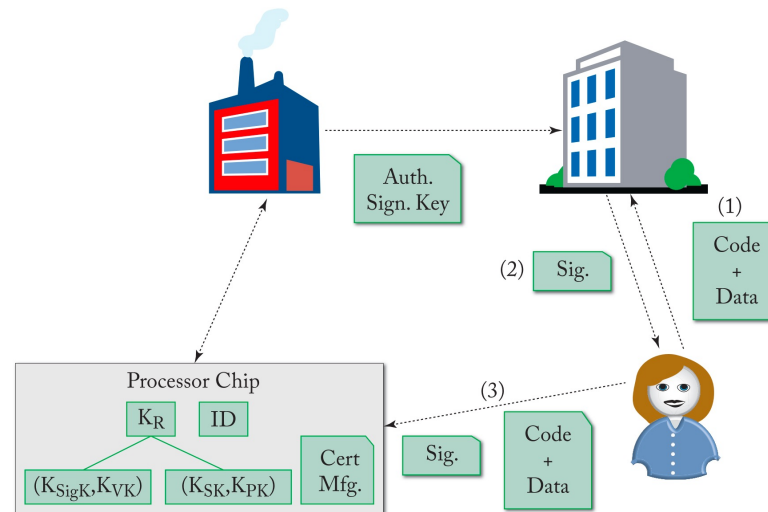


Butterfly PUF:

Excitation in preset and clr creates instability in cross coupled inputs but settles in a stable output

Limiting the Execution to Only Authorized Code

- TCB updates (or protected) software can be authenticated with a separate key (manufacturer dependent)
- Add **public key** to the processor at manufacture time
- Only software signed by manufacturer can be sent to the processor
- Code can be **signed under demand** using also processor keys (e.g., to prevent install certain firmware out of a particular time window)



Lock-In and Privacy Concerns

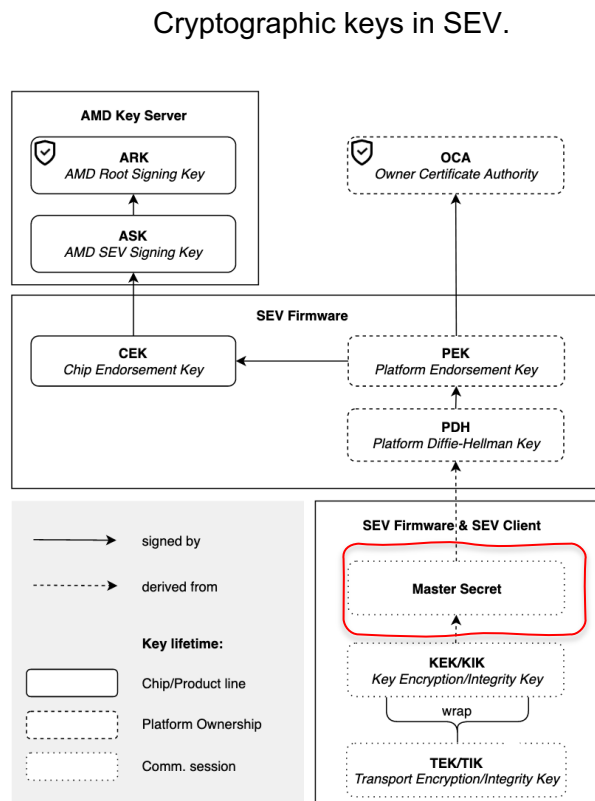
- ▣ User depends on manufacturer to send information to TEE
 - ◆ E.g., Software running inside SGX requires **Intel signature**
- ▣ When TCB and TEE is authenticated, use processor specific keys that can be used by the Certification Authority to infer information about the systems requesting the update
- ▣ **Direct Anonymous Attestation** (DAA) on TPM addresses these issues (partially)
 - ◆ Camenisch-Lysyanskaya Signatures

Root of Trust Assumptions

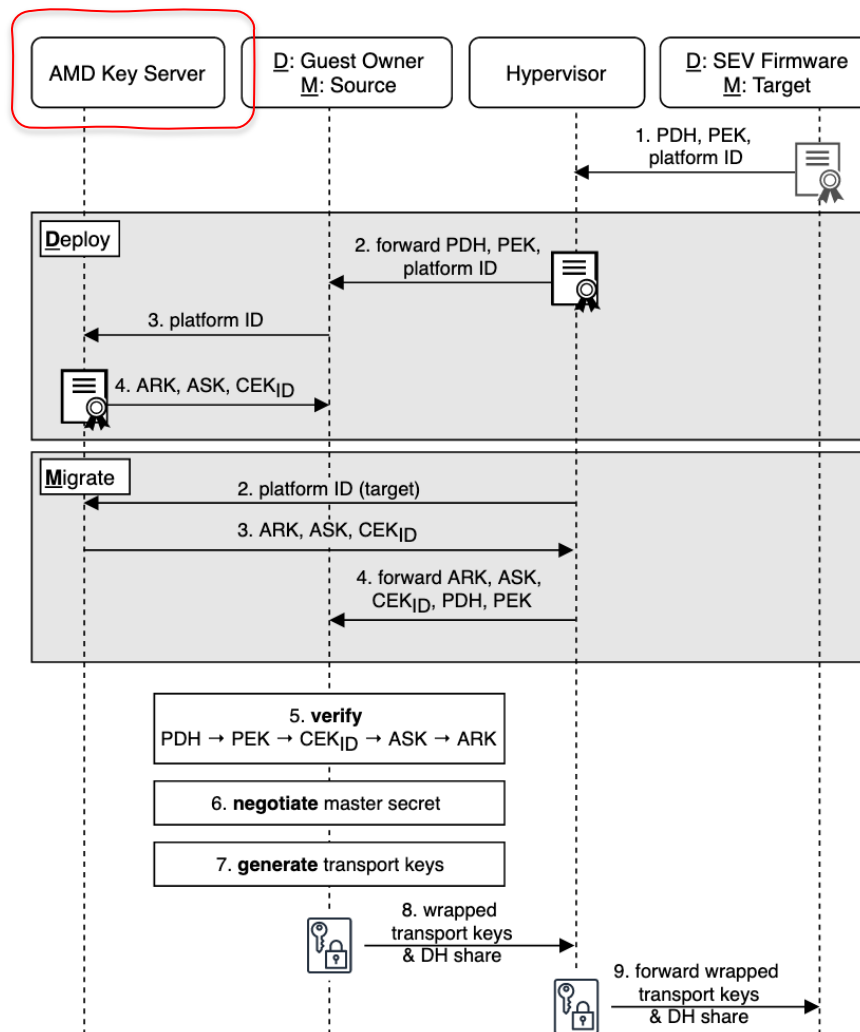
- ▣ Unique key assumptions
 - ◆ Manufacturer can never reuse the same key or use PUF
- ▣ Protected root of trust assumption
 - ◆ **Kr will never be disclosed**
 - ◆ TCB never will leak it
- ▣ Fresh measurement assumption
 - ◆ Need to be fresh to prevent TOC-TOU attacks

Example Attack : Remote Attestation with AMD SEV

■ <https://arxiv.org/pdf/1908.11680.pdf>



SEV Secure Chanel (M: migration, D: deploy)



Example Attack: AMD TCB exploit for attestation

- ❑ **One Glitch to Rule Them All (2021)** <https://arxiv.org/abs/2108.04575>
 - ◆ <https://github.com/PSPReverse/amd-sev-migration-attack>
 - ◆ Voltage glitching attack that **allows an attacker to execute custom payloads on the AMD-SPs** --> Extract keys to deploy "custom" firmware
 - Allow to access secured VM using live-migration
 - ◆ Allows to generate **false attestation reports** (in the "custom" firmware server) (because we can change root of trust and fake the secure protocol)
 - An attacker without physical access to the target host can use extracted endorsement keys to attack SEV-protected VM s

Example Attack : SGX Attacks

■ <https://arxiv.org/pdf/2006.13598.pdf>

Attacks (abbreviated titles)	Type	SGX Specific	Targeted attack	Impact						Mitigations			
				Page access pattern	Instruction trace	Instruction latency	Memory access pattern	Memory Contents	Fault Injection	Microcode patch	System design	Compiler/SDK	Application design
Controlled-Channel [9]	sec. III-A	●	●	●	○	○	○	○	○	○	○	● ^[52]	● ^[55]
Stealthy Page Table [10]	sec. III-A	●	●	●	○	○	○	○	○	○	○	● ^[52]	○
SGX-Step [11]	sec. III-A	●	○	●	●	●	○	○	○	○	○	● ^[11] ● ^[48]	○
Nemesis [12]	sec. III-A	●	●	○	○	●	○	○	○	○	○	● ^[50]	○
Off Limits [13]	sec. III-A	●	●	●	●	○	○	○	○	● ^[13]	○	○	● ^[13]
Leaky Cauldron [14]	sec. III-B	●	●	●	○	○	●	○	○	○	○	● ^[50] ● ^[51]	○
CacheZoom [20]	sec. III-B	●	●	○	○	○	●	○	○	○	○	● ^[50] ● ^[51]	● ^[56]
Cache Attacks on SGX [21]	sec. III-B	●	●	○	○	○	●	○	○	○	○	○	● ^[56]
Malware Guard Extensions [22]	sec. III-B	●	●	○	○	○	●	○	○	○	● ^[22]	○	○
Software Grand Exposure [23]	sec. III-B	●	●	○	○	○	●	○	○	○	○	○	● ^[23]
CacheQuote [24]	sec. III-B	●	●	○	○	○	●	○	○	○	● ^[24]	○	○
MemJam [25]	sec. III-B	○	●	○	○	○	●	○	○	○	○	● ^[49]	○
Branch Shadowing [26]	sec. III-C	●	●	○	●	○	○	○	○	○	○	● ^[26] ● ^[48]	○
BranchScope [27]	sec. III-C	○	●	○	●	○	○	○	○	○	○	● ^[27] ● ^[48]	○
Bluethunder [28]	sec. III-C	●	●	○	●	○	○	○	○	○	○	● ^[28]	○
SgxPectre [31]	sec. III-D	●	●	○	○	○	○	●	○	● ^[47]	○	○	○
SpectreRSB [32]	sec. III-D	○	●	○	○	○	○	●	○	○	○	○	● ^[32]
Spectre v1 [33]	sec. III-D	○	●	○	○	○	○	●	○	● ^[47]	○	○	○
Foreshadow-SGX [34]	sec. III-E	●	○	○	○	○	○	●	○	● ^[35]	○	○	○
RIDL [38]	sec. III-F	○	●	○	○	○	○	●	○	● ^[38]	○	○	○
ZombieLoad [39]	sec. III-F	○	●	○	○	○	○	●	○	● ^[39]	○	○	○
CacheOut [40]	sec. III-F	●	●	○	○	○	○	●	○	● ^[40]	○	○	○
CrossTalk [41]	sec. III-F	●	○	○	○	○	○	●	○	● ^[41]	○	○	○
Plundervolt [44]	sec. III-G	●	○	○	○	○	○	○	●	● ^[44]	○	○	○