

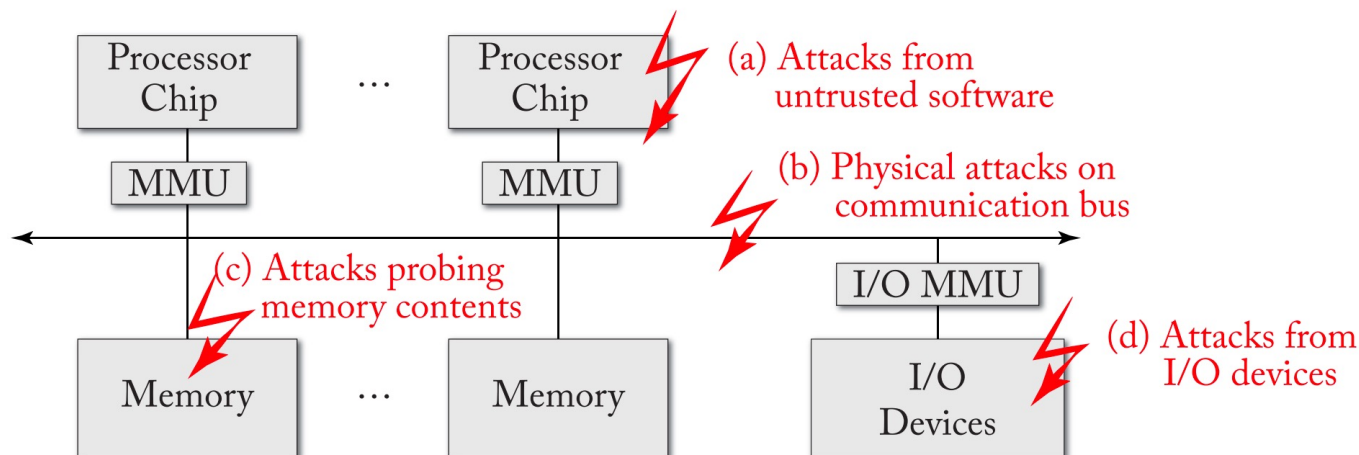
Memory Protection

Chapter 6

- [1] J. Szefer, “Principles of secure processor architecture design,” *Synth. Lect. Comput. Archit.*, vol. 13, no. 3, pp. 1–173, 2018.

Threats Against Main Memory

- Wiring and memory device itself are untrusted assumed built upon DRAM devices
- Sources of Attacks
 - ◆ Untrusted software (bypassing isolation barriers)
 - ◆ Malicious devices (trying to access via DMA to protected regions)
 - ◆ Physical attacks (on memory "bus" via probing)
 - ◆ Physical attacks (on memory itself, Cooldboot, Rowhammer)
- Memory can be easily removed and analyzed off-line
- Competing Non-volatile Memories (NVM) in the horizon (such as Intel Optane) can make this problem harder



Attacks to Memory

▣ Passive Attacks

- ◆ **Eavesdropping attacks:** observe information or accessing patterns without altering it to gain knowledge
 - Pattern: e.g., Observing AES S-Box access pattern can leak information about the encryption key
 - Prevented: **encryption** and **obfuscation**

▣ Active Attacks


- ◆ **Spoofing attacks:** inject memory data (or operations) without being detected
 - Inject changes in memory mapping, processor instructions, stole data, ...
 - Prevented via **hashing**
- ◆ **Splicing Attacks:** combine multiple read/write operations in a new (legitimate) read or write
 - Splice parts of different messages (e.g., payload from one and header from other)
 - Prevented via **keyed hashing** (MAC Message Authenticated Code)
- ◆ **Replay Attacks:** Send messages again
 - Reuse old "known" messages (v.gr. replace an encryption key)
 - Prevented via **nonced hashing**

Main Memory Protection Mechanisms

- ▣ Three main techniques and objectives
 - ◆ **Confidentiality**: with encryption
 - ◆ **Integrity**: with hashing (typically hash-trees)
 - ◆ **Access Pattern protection** with access pattern obfuscation
- ▣ Memory protection is focused on off-chip related issues
 - ◆ Encryption and/or hashing is only done at **off-chip processor interfaces**
 - ◆ Inside the processor chip regular isolation techniques (e.g., page tables for equally privileged software or tagging for higher privilege software)
 - ◆ Might increase average access time
 - ◆ Access pattern still discoverable
- ▣ MMU and IOMMU are also useful

Confidentiality Protection With Encryption

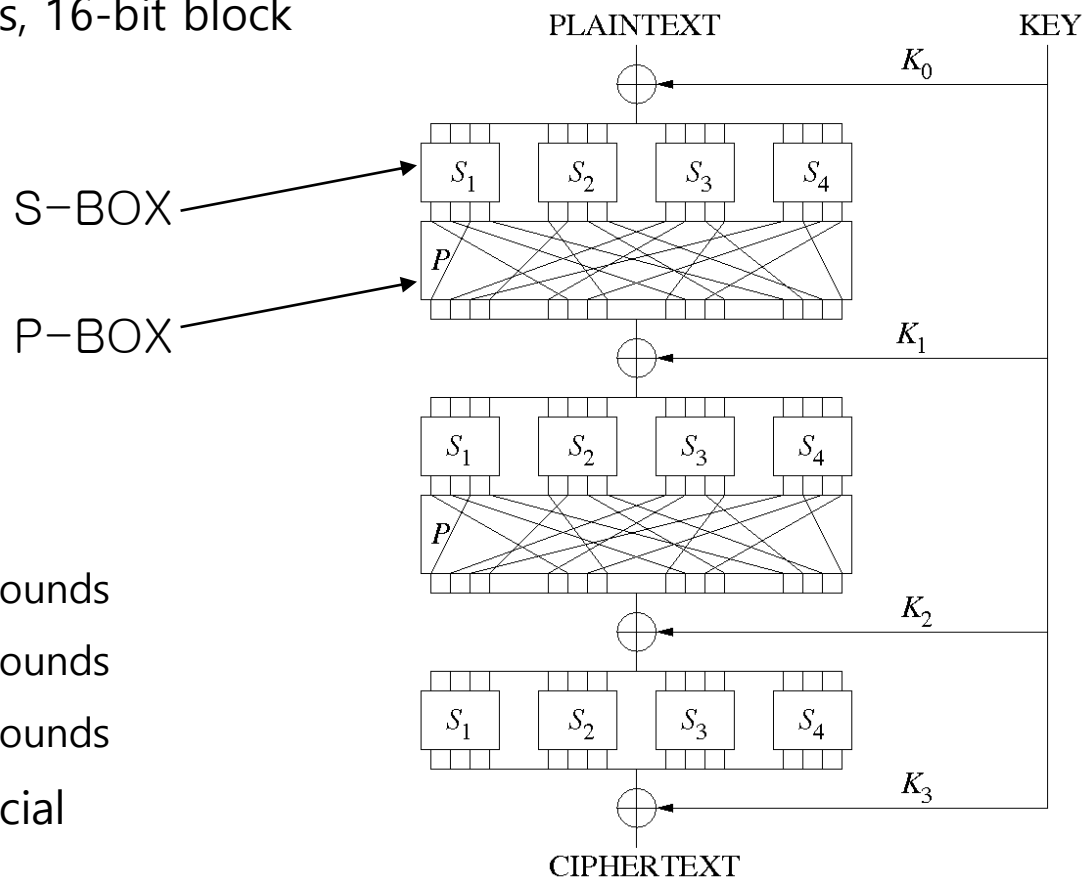
- Use a **re-generated** key each time system **reboot**
- Store the key in the MMUs or special management module (usually Memory Controller)
- Application, OS or system level. Can be **selectively** enabled in some ranges of the addressing space of the process or physical addresses of the memory (use ASIDs in tags to isolate)
- Use asymmetric cryptography to support multi-socket systems (the key never should leave the processor chip in plaintext!!) or use independent key per memory controller/ASID
- Memory access will be slower for encrypted data (use **suitable algorithms** for the task: AES CTR Mode)
- To support DMA IOMMU also should have a key (I/O and PCIe is unencrypted, therefore vulnerable against probing)
- Can be vulnerable if management engine is untrusted (i.e., having access to encryption keys)
- AMD Secure Encrypted Virtualization implements it



Aside: AES Algorithm (10.000 feet view)

- Base on substitution-permutation network

- Example 3-rounds, 16-bit block



- AES

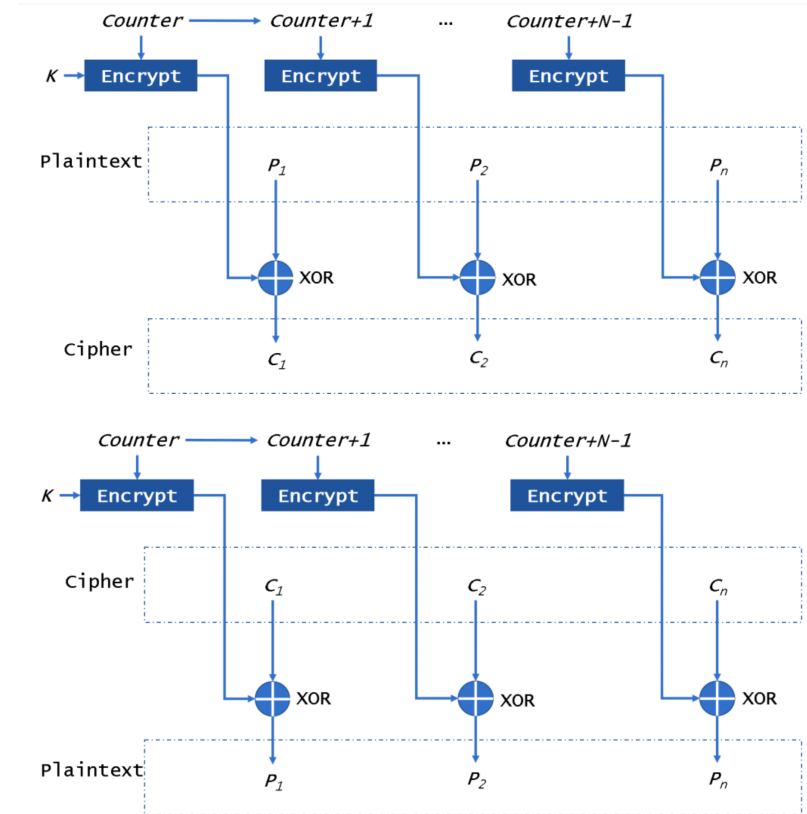
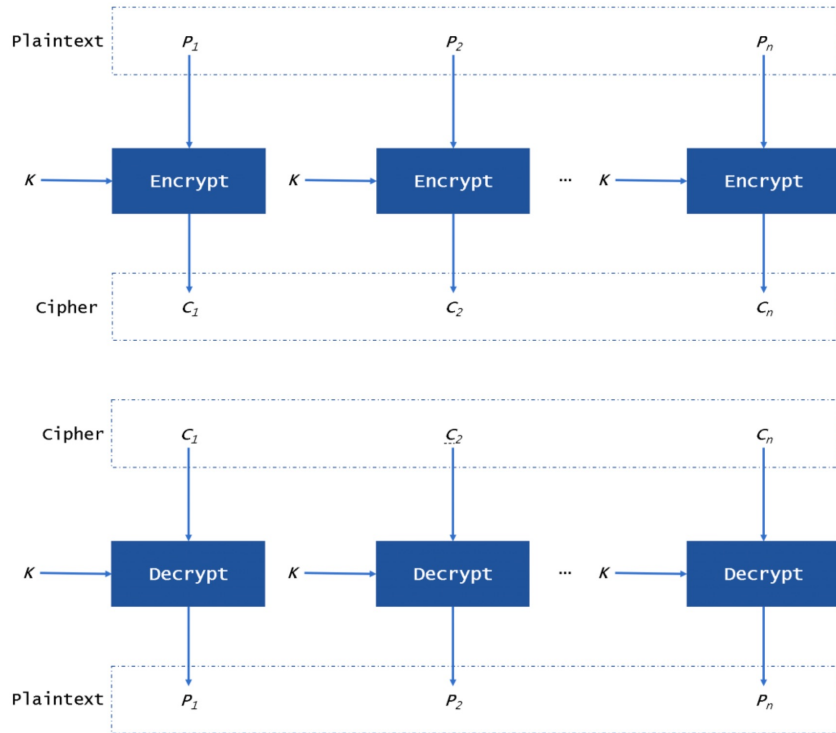
- 128-bit blocks
 - 128bit key: 10rounds
 - 256bit key: 14rounds
 - 192bit key: 12rounds
- Last round is special

- AES-NI

- Instructions for one round of ENC/DEC

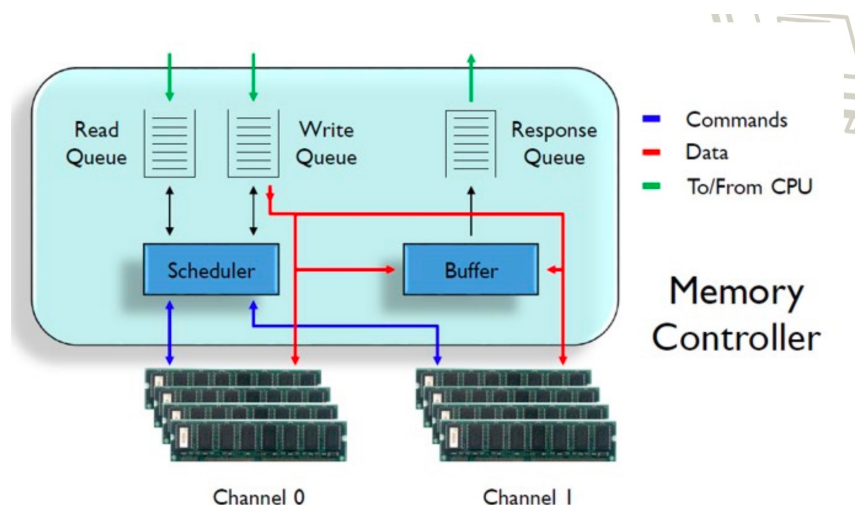
Aside: Counter Mode AES (fast encryption/decryption)

■ Electronic Code Book vs Counter Mode (AES has 5 modes)



Aside: Memory Controller and Memory Encryption Engine

- ❑ Quite complex piece (as much as a processor)
- ❑ Connects the on-chip network with the DRAM
- ❑ Handles
 - ◆ DRAM refresh
 - ◆ Row-buffer management policies
 - ◆ Address mapping Schemes
 - ◆ Request scheduling
- ❑ Now also encryption and decryption
 - ◆ Memory Encryption Engine (MEE)



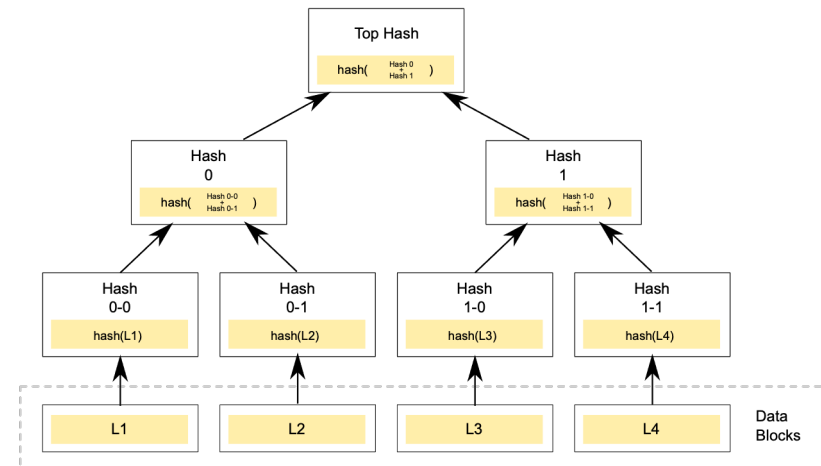
Examples of MEE

- ▣ AMD SEV (and SME) modes
 - ◆ 128bit AES XEX (?) (xor-encrypt-xor) [AMD Rome]
 - ◆ More resistant to counter reuse than CTR
 - ◆ We will see...

- ▣ Intel SGX (Skylake) <https://eprint.iacr.org/2016/204.pdf>
 - ◆ CTR (?)
 - ◆ 15 cycles for a write, 1block per cycle (pipeline)
 - ◆ 5.5% performance degradation in gobmk (SPECInt 2006)
 - ◆ Includes also integrity

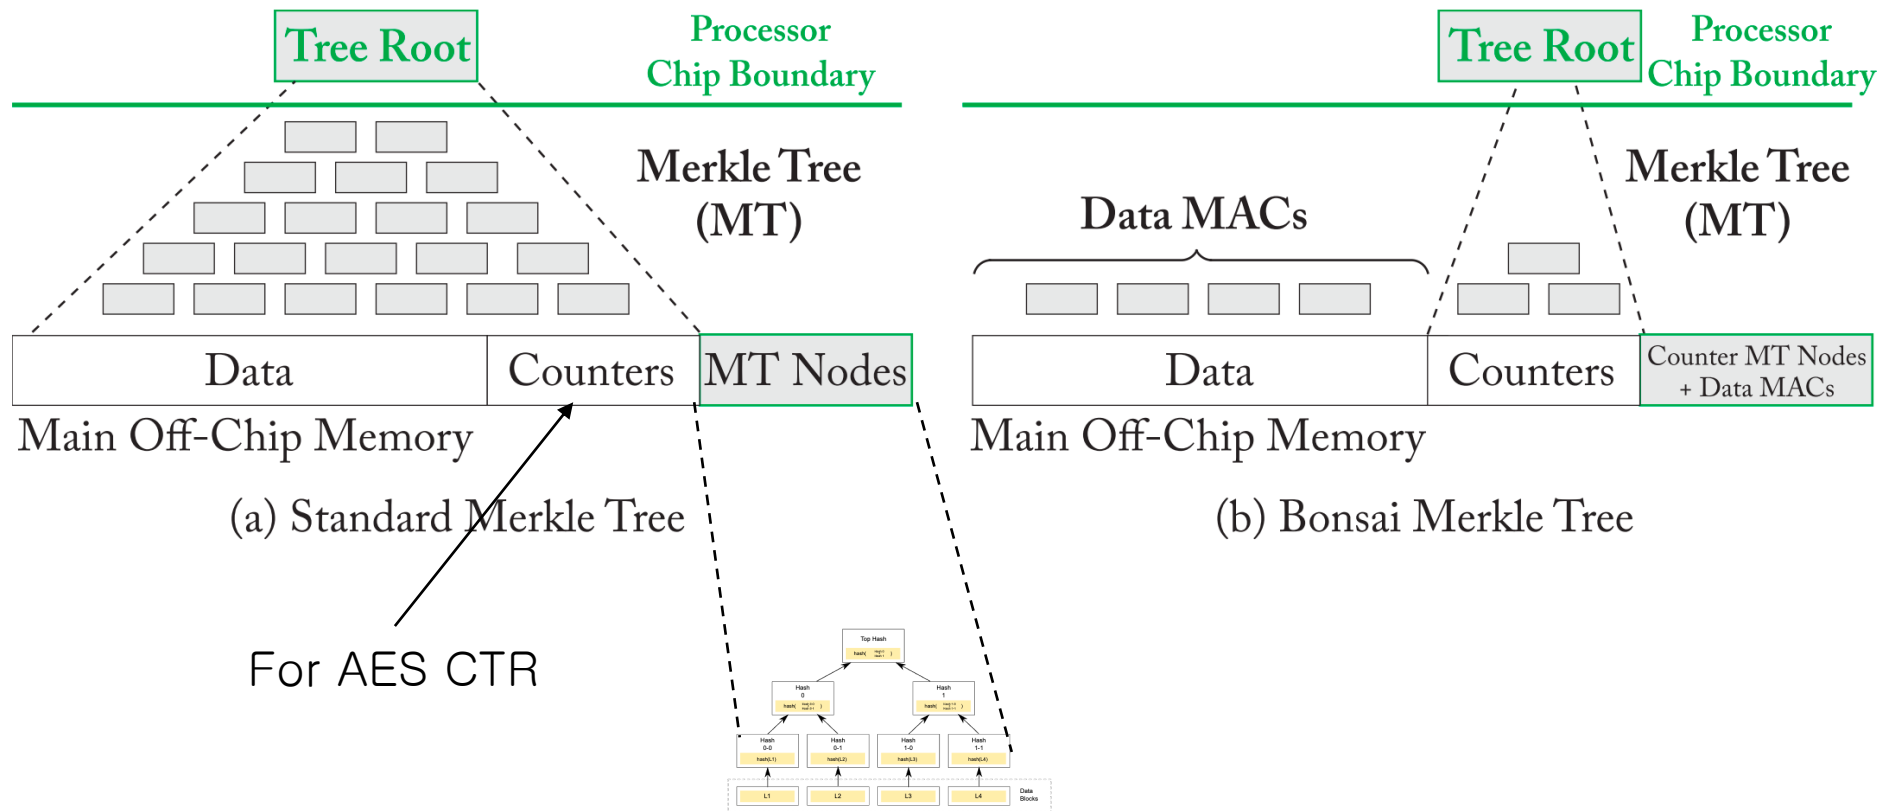
Integrity Protection With Hashing

- ❑ Cryptographic hashes
 - ◆ Produces a unique output for an input (negligible probability of collision of two inputs if the hash is long enough)
 - ◆ We can't figure out the input from the output (non inversible functions)
- ❑ The result of hashing the data can be understood as a fingerprint of the data
 - ◆ The easiest way to check if the data has been altered is check hash result
- ❑ Integrity protections focused on **external attackers**
- ❑ Hash the whole memory is **impractical**: slow and every change requires to recompute the hash: **Use hash trees**
- ❑ $\log_2(N)$ hash computations (N: depth of the tree)
 - ◆ **Write**: update all the hashes from the leaves updated to the root
 - ◆ **Read**: check if involved hashes to the data are ok
- ❑ But attacker can modify the whole tree (??)
 - ◆ Add encryption to the hash (MAC)



Merkle Trees

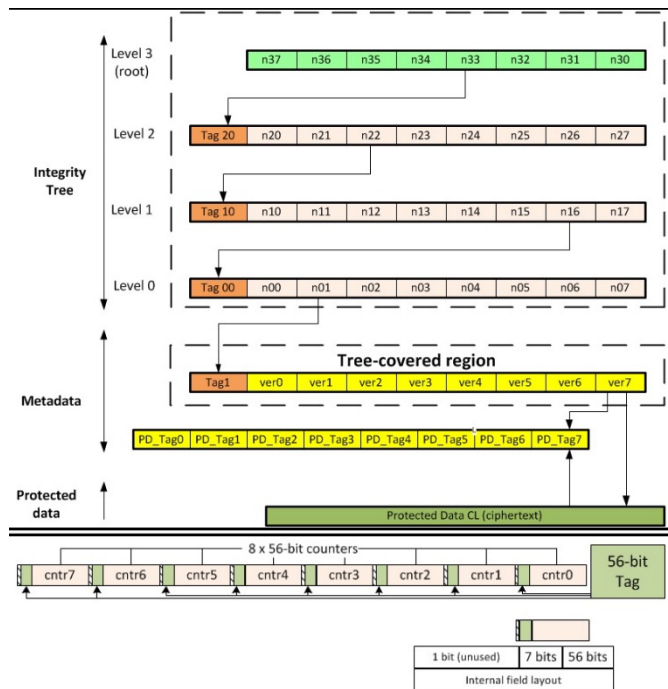
- Root of tree on secure on-chip register



Hash Tree Protection

- Protect intermediate nodes in the tree: attacker can compute and insert its own hashes,
 - ◆ **Do not use plain hashes**: use secure processor-based keys to encrypt the hash (the tree is in memory!)
 - ◆ Use cryptographic hashes: **keyed hashes** (**MAC** or Message Authenticate Code) and always on secure on-chip locations: can't be tampered (but can be cached)
- MAC can protect against splicing and spoofing but **not replaying**
 - ◆ Add a monotonic counter for **replay attacks** protection
- **Secure page swapping** might require to reserve leaf nodes in the hash tree for non-present pages (i.e., integrity protection trees should cover the swapped pages)
- Variety of **performance improvement** to minimize hashing performance impact
 - ◆ Counter mode encryption (out of the critical path)
 - ◆ Bonsais Merkle independent encryption seed from the memory address (do not combine key with address in the original tree)
- Intel SGX uses a tree of MACs for integrity and replay protection
 - ◆ <https://eprint.iacr.org/2016/204.pdf>

SGX Bonsai Merkle Tree



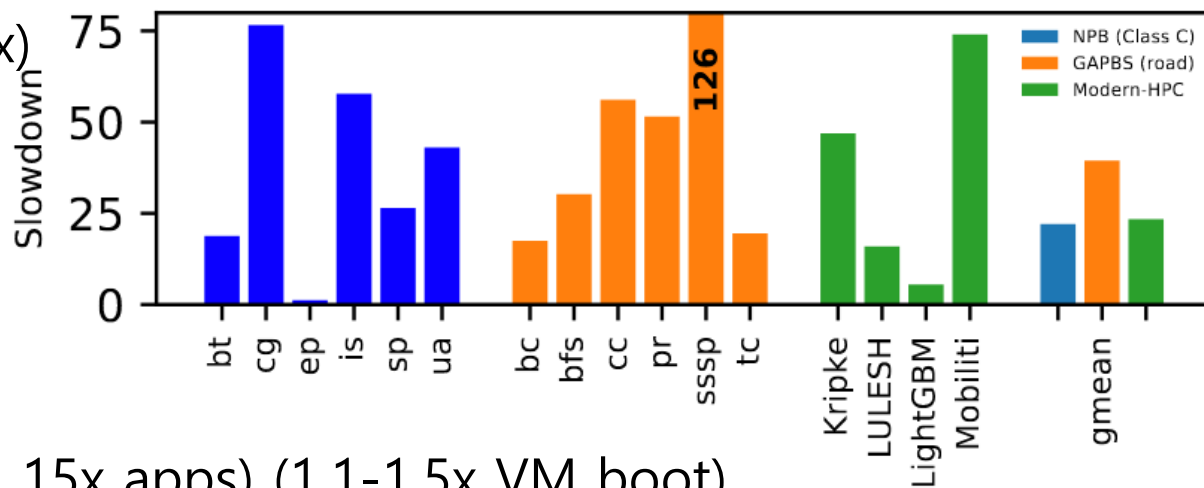
	Start offset	End offset	Region size
Protected data	000_0000	5ff_ffff	96MB
Version + PD_Tag	600_0000	77f_ffff	24MB
Reserved	780_0000	7df_ffff	6MB
L0	7e0_0000	7f7_ffff	1.5MB
Reserved	7f8_0000	7fb_ffff	256KB
L1	7fc_0000	7fe_ffff	192K
Reserved	7ff_0000	7ff_7fff	32KB
L2	7ff_8000	7ff_dfff	24K
Reserved	7ff_e000	7ff_ffff	4KB
L3 (On-die root)	7ff_f000	7ff_ffff	4KB
Total MEE region			128MB

Performance impact?

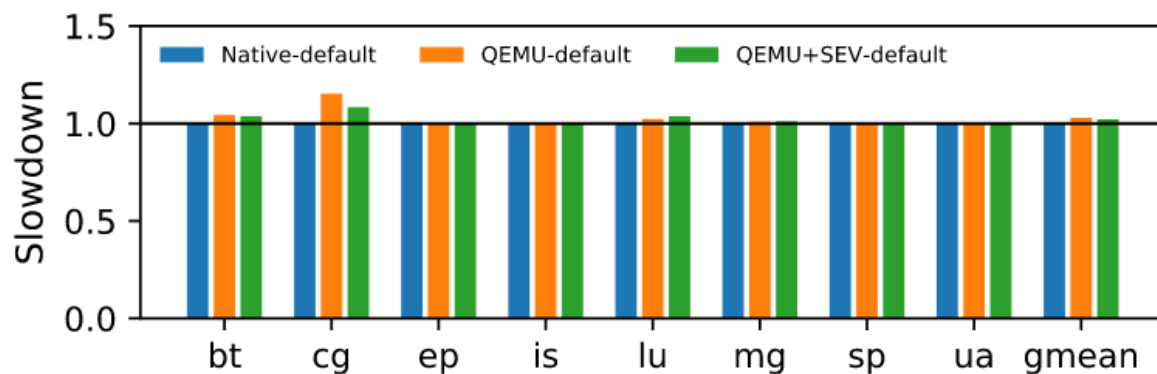
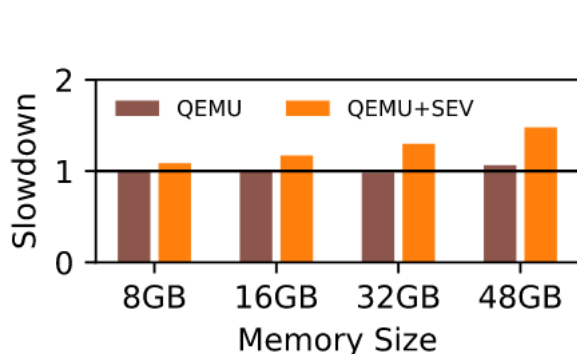
- HPC Apps from <https://arxiv.org/pdf/2010.13216.pdf>

- ◆ Single node!

- SGX (1.2x-126x)

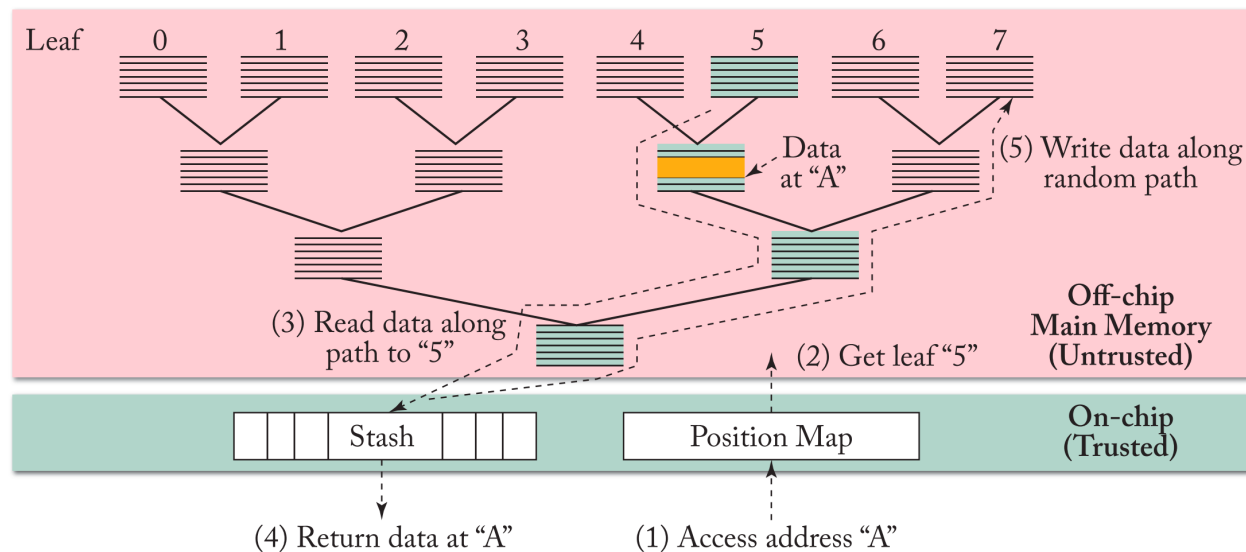


- SEV/SME (1x-1.15x apps) (1.1-1.5x VM boot)



Access Pattern Protection

- Just observing the access to address activity can be used as conduit for revealing secrets
- ORAM (Oblivious RAM) keeps the semantics of the program but hides the access pattern by **shuffling memory locations**
 - On each access memory insert random accesses (to dummy data). The objective is to make indistinguishable to the attacker the useful accesses from the noise
 - Only in sensible data (software + hardware, e.g., S-Box access in AES encryption)



- All OS does something called Address Space Layout Randomization (ASLR) but oriented to prevent exploits (e.g., buffer overflows). Its at page level.